

Machine Learning Classification Algorithms to Car Evaluation Dataset

Team Members-

Eshanth Patyal(21ucs078)

Harsh Gupta(21ucs084)

Harsh Agarwal(21ucs083)

Yash R Khandelwal(21ucs254)

Contents

1. Objective
2. Specifications of the Dataset
3. Dataset Description
4. Libraries used
5. Statistical summary
6. Data Preprocessing
 - 6.1. Labelling Data
7. Implementing the algorithms
 - 7.1. SVM
 - 7.2. K-Nearest-Neighbours (KNN)
 - 7.3. Naive Bayes
 - 7.4. Logistic Regression
 - 7.5. Decision Tree
 - 7.6. Random Forest Classifier
8. Conclusion

1.Objective

The project aims to perform preprocessing on a dataset and realise the Different machine learning algorithm learned in the classroom to the data. The dataset contains 6 dependent variables and 1 target variable.

2.Specifications of the Dataset

The dataset is taken from the UCI machine learning repository.

Dataset Characteristics-Multivariate

Associated Tasks-Classification

Feature Type -Categorical

Number of Instances 1728

Number of Features 6

Class Labels-

Unacc(unaccepetable), acc(accapetable), good, vgood

Variable Name	Role	Type	Description	Missing Values	Variable Name
buying	Feature	Categorical	buying price	no	buying
maint	Feature	Categorical	price of the maintenance	no	maint
doors	Feature	Categorical	number of doors	no	doors
persons	Feature	Categorical	capacity in terms of persons to carry	no	persons
lug_boot	Feature	Categorical	the size of luggage boot	no	lug_boot
safety	Feature	Categorical	estimated safety of the car	no	safety

3.Dataset Description

The Dataset is split into to dataframes-

Df_x=the 6 features

Df_y=the target variable

```
[5] 1
    2 print(car_evaluation.variables)
```

	name	role	type	demographic \
0	buying	Feature	Categorical	None
1	maint	Feature	Categorical	None
2	doors	Feature	Categorical	None
3	persons	Feature	Categorical	None
4	lug_boot	Feature	Categorical	None
5	safety	Feature	Categorical	None
6	class	Target	Categorical	None

		description	units	missing_values
0		buying price	None	no
1		price of the maintenance	None	no
2		number of doors	None	no
3		capacity in terms of persons to carry	None	no
4		the size of luggage boot	None	no
5		estimated safety of the car	None	no
6		evaluation level (unacceptable, acceptable, go...	None	no

4.Libraries Used

```
[1] 1 pip install ucimlrepo
```

Collecting ucimlrepo
Downloading ucimlrepo-0.0.3-py3-none-any.whl (7.0 kB)
Installing collected packages: ucimlrepo
Successfully installed ucimlrepo-0.0.3

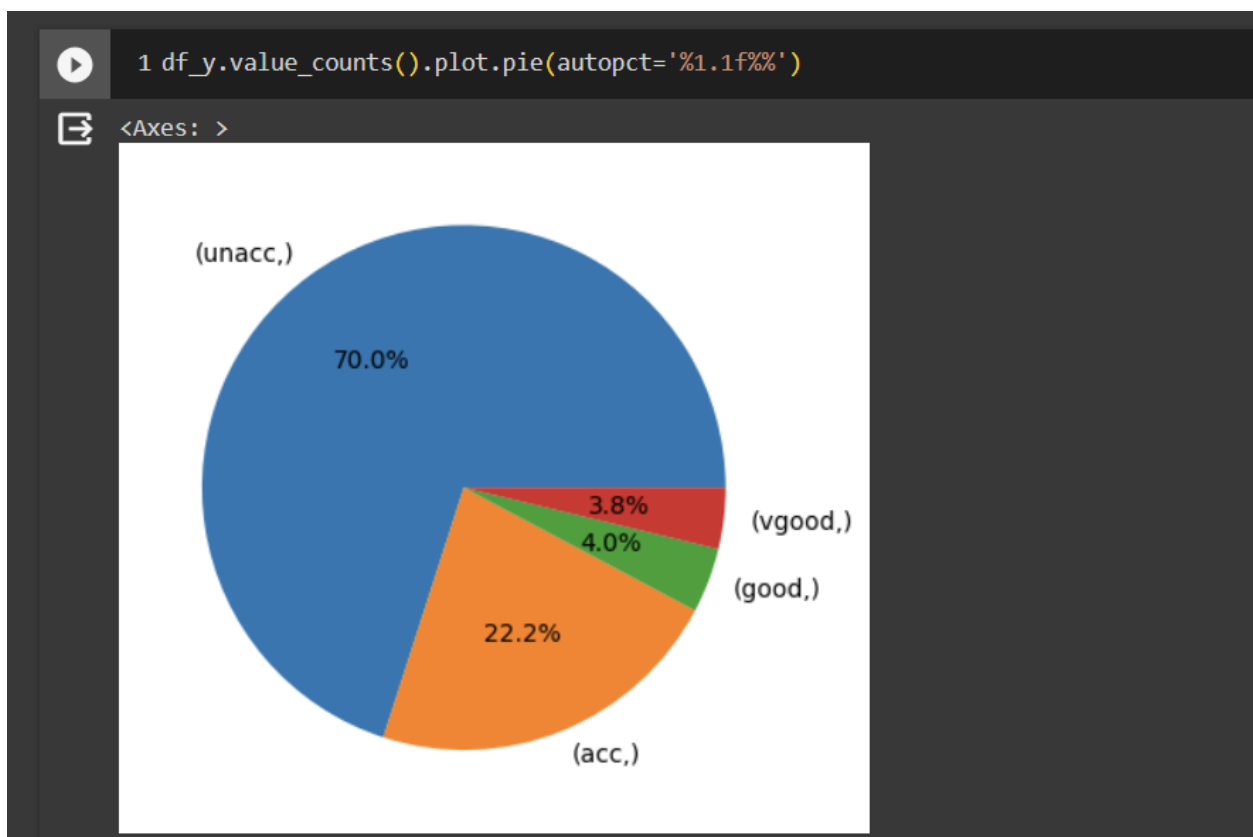
```
1 import pandas as pd
2 import seaborn as sns
3 import matplotlib.pyplot as plt
4 import numpy as np
5 import sklearn
```

UCI ML repo imported for the dataset.


```
✓ 0s 1 from ucimlrepo import fetch_ucirepo
2
3 # fetch dataset
4 car_evaluation = fetch_ucirepo(id=19)
5
6 # data (as pandas dataframes)
7 df_x = car_evaluation.data.features
8 df_y = car_evaluation.data.targets
9
10 # metadata
11 print(car_evaluation.metadata)
12
```

```
{'uci_id': 19, 'name': 'Car Evaluation', 'repositor
```

5. Statistical Summary



Introduction To Data Science Project



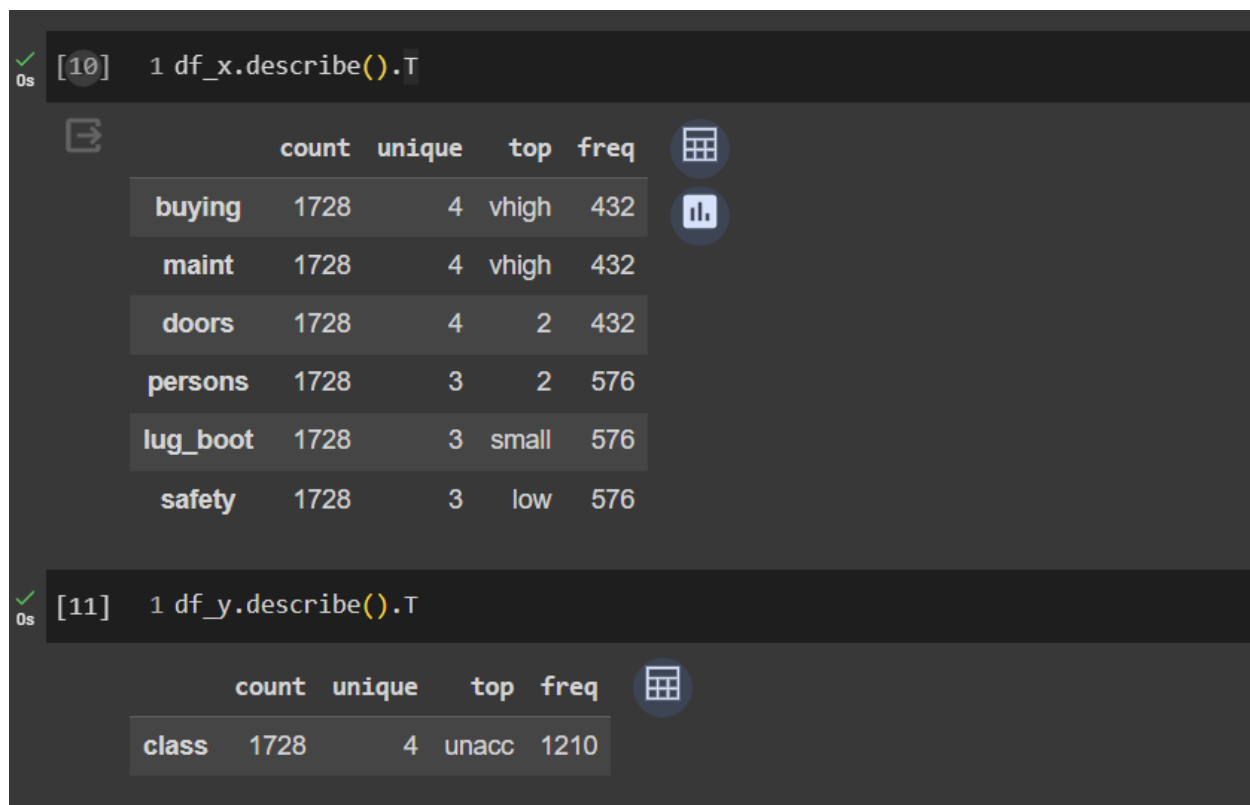
- Representation of data in class labels is more in unacc class.
- Distribution of all independent variables is evenly spread.

```
1 df_x.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1728 entries, 0 to 1727
Data columns (total 6 columns):
Column Non-Null Count Dtype
--- ---
0 buying 1728 non-null object
1 maint 1728 non-null object
2 doors 1728 non-null object
3 persons 1728 non-null object
4 lug_boot 1728 non-null object
5 safety 1728 non-null object
dtypes: object(6)
memory usage: 81.1+ KB

```
[8] 1 df_y.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1728 entries, 0 to 1727
Data columns (total 1 columns):
Column Non-Null Count Dtype
--- ---
0 class 1728 non-null object
dtypes: object(1)
memory usage: 13.6+ KB



```
[10] 1 df_x.describe().T
```

	count	unique	top	freq
buying	1728	4	vhigh	432
maint	1728	4	vhigh	432
doors	1728	4	2	432
persons	1728	3	2	576
lug_boot	1728	3	small	576
safety	1728	3	low	576

```
[11] 1 df_y.describe().T
```

	count	unique	top	freq
class	1728	4	unacc	1210

6. Data Preprocessing

Data preprocessing is a technique for data mining that involves the translation of raw data into a comprehensible format. In this process, the information used was pre-processed. Since the dataset didn't have any empty values, there was no need to change NaN values.

The issue was Non-Numerical values in many attributes which needed to be changed.

6.1 Labelling Data(one hot coding)

We will see that each record has several non-numeric attributes, such as maint, buying, etc., from the available data collection. Generally, learning algorithms need numerical data.

We use the one-hot encoding method, one common way to convert categorical variables to numerical variables.


```
1 from sklearn.preprocessing import OneHotEncoder
2 import pandas as pd
3
4 # Concatenate df_x and df_y to perform one-hot encoding on the entire dataset
5 df = pd.concat([df_x, df_y], axis=1)
6
7 # Perform one-hot encoding using pd.get_dummies
8 df_encoded = pd.get_dummies(df, columns=['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety'])
9 df_class=pd.get_dummies(df_y)
10
11 df_x_encoded = df_encoded.drop(['class'], axis=1)
12 df_y_encoded = df_class
```

After encoding data looks like-

```
1 df_x_encoded[0:5]
```

	buying_high	buying_low	buying_med	buying_vhigh	maint_high	maint_low	maint_med	maint_vhigh	doors_2	doors_3	...	doors_5more	persons_2	persons_4	persons_more	lug_boot_big
0	0	0	0	1	0	0	0	1	1	0	...	0	1	0	0	0
1	0	0	0	1	0	0	0	1	1	0	...	0	1	0	0	0
2	0	0	0	1	0	0	0	1	1	0	...	0	1	0	0	0
3	0	0	0	1	0	0	0	1	1	0	...	0	1	0	0	0
4	0	0	0	1	0	0	0	1	1	0	...	0	1	0	0	0

5 rows x 21 columns

```
[14] 1 df_y_encoded[0:5]
```

	class_acc	class_good	class_unacc	class_vgood
0	0	0	1	0
1	0	0	1	0
2	0	0	1	0
3	0	0	1	0
4	0	0	1	0

7.implementing the algorithm

7.1 SVM

SVM is effective when dealing with many different decision features feature, making it suitable for datasets with multiple independent variables.

SVM can handle both linear and non-linear decision planes. In this case, where the relationship between features and the target variable is not linear (6 independent variables which are all different), SVM can capture complex patterns.

SVM is robust and is less prone to overfitting in such non-linear relations.

```
1 from sklearn.svm import SVC
2
3
4 svc_classifier = SVC()
5
6 svc_classifier.fit(X_train, y_train)
7 y_pred_svc = svc_classifier.predict(X_test)
8
9 # Evaluate the model
10 accuracy_svc = accuracy_score(y_test, y_pred_svc)
11 classification_report_svc = classification_report(y_test, y_pred_svc)
12
13 print(f"SVC Accuracy: {accuracy_svc:.2f}")
14 print("SVC Classification Report:\n", classification_report_svc)
15
```

SVC Accuracy: 0.97
SVC Classification Report:

	precision	recall	f1-score	support
acc	0.99	0.89	0.94	83
good	0.59	0.91	0.71	11
unacc	1.00	1.00	1.00	235
vgood	0.84	0.94	0.89	17
accuracy			0.97	346
macro avg	0.85	0.94	0.88	346
weighted avg	0.98	0.97	0.97	346

SVM has provided with exceptional accuracy in unacc class which has dominated the dataset(70% of the class labels). The same is true for acc and vgood class. Only good class label (4% of class labels)has low precision.

7.2 KNN

KNN is a simple and intuitive algorithm that can be effective when the decision boundaries are not well-defined. It works well when there are regions in the feature space where certain classes are concentrated.

```
[20] 1 y = df_y
      2 X_train, X_test, y_train, y_test = train_test_split(df_x_encoded, y, test_size=0.2, random_state=42)
      3
```

```
[21] 1 from sklearn.neighbors import KNeighborsClassifier
      2
      3 # Instantiate the k-Nearest Neighbors Classifier
      4 knn_classifier = KNeighborsClassifier()
      5
      6 # Train the classifier
      7 knn_classifier.fit(X_train, y_train)
      8
      9 # Make predictions on the test set
     10 y_pred_knn = knn_classifier.predict(X_test)
     11
     12 # Evaluate the model
     13 accuracy_knn = accuracy_score(y_test, y_pred_knn)
     14 classification_report_knn = classification_report(y_test, y_pred_knn)
     15
     16 print(f"KNN Accuracy: {accuracy_knn:.2f}")
     17 print("KNN Classification Report:\n", classification_report_knn)
     18
```

KNN Accuracy: 0.88
KNN Classification Report:

	precision	recall	f1-score	support
acc	0.78	0.75	0.77	83
good	0.30	0.27	0.29	11
unacc	0.92	0.99	0.96	235
vgood	1.00	0.29	0.45	17
accuracy			0.88	346
macro avg	0.75	0.58	0.62	346
weighted avg	0.87	0.88	0.86	346

Almost a similar result is seen as good class label is still showing low overall correctness, whereas all other classes seem to do much better in comparison.

7.3 Naïve Bayes

It performs well with categorical data and assumes independence between features, making it a good choice when dealing with categorical variables like 'buying', 'maint', etc.

```
Naive Bayes

[23] 1 from sklearn.naive_bayes import MultinomialNB
      2
      3 nb_classifier = MultinomialNB()
      4 nb_classifier.fit(X_train, y_train)
      5 y_pred_nb = nb_classifier.predict(X_test)
      6
      7 # Evaluate the model
      8 accuracy_nb = accuracy_score(y_test, y_pred_nb)
      9 classification_report_nb = classification_report(y_test, y_pred_nb)
     10
     11 print(f"Naive Bayes Accuracy: {accuracy_nb:.2f}")
     12 print("Naive Bayes Classification Report:\n", classification_report_nb)
```

Naive Bayes Accuracy: 0.82

Naive Bayes Classification Report:

	precision	recall	f1-score	support
acc	0.63	0.54	0.58	83
good	0.57	0.36	0.44	11
unacc	0.87	0.97	0.91	235
vgood	1.00	0.35	0.52	17
accuracy			0.82	346
macro avg	0.77	0.56	0.62	346
weighted avg	0.81	0.82	0.80	346

Results from naïve bayes seem to be more like KNN than SVM but the performance of the model is downgraded from KNN.


7.4 Logistic Regression

Logistic Regression is easy to interpret and provides probabilities for class assignments. Logistic Regression assumes a linear relationship between the

independent variables and the log-odds of the target variable, making it suitable for datasets where such assumptions hold.

Logistic Regression

```
1 from sklearn.linear_model import LogisticRegression
2
3 logreg_model = LogisticRegression()
4 logreg_model.fit(X_train, y_train)
5 y_pred_logreg = logreg_model.predict(X_test)
6
7 # Evaluate the model
8 accuracy_logreg = accuracy_score(y_test, y_pred_logreg)
9 classification_report_logreg = classification_report(y_test, y_pred_logreg)
10
11 print(f"Logistic Regression Accuracy: {accuracy_logreg:.2f}")
12 print("Logistic Regression Classification Report:\n", classification_report_logreg)
13
```

 Logistic Regression Accuracy: 0.92
Logistic Regression Classification Report:

	precision	recall	f1-score	support
acc	0.84	0.82	0.83	83
good	0.50	0.55	0.52	11
unacc	0.96	0.97	0.97	235
vgood	0.94	0.88	0.91	17
accuracy			0.92	346
macro avg	0.81	0.80	0.81	346
weighted avg	0.92	0.92	0.92	346

Logistic regression model performs better than KNN and Naïve bayes but not as good as SVM.

The same issues are still consistent (good class label showing low performance) as with other models.

7.5 Decision Tree

Decision trees are interpretable and easy to understand, providing insights into feature importance. They can handle both numerical and categorical

data without the need for lots of pre-processing. Decision trees can capture non-linear relationships and interactions between features, making them suitable for this dataset.

Decision Tree

```
[25] 1 from sklearn.tree import DecisionTreeClassifier
      2
      3 # Init of Decision Tree
      4 dt_classifier = DecisionTreeClassifier()
      5
      6 # Train the classifier
      7 dt_classifier.fit(X_train, y_train)
      8
      9 # Make predictions on the test set
     10 y_pred_dt = dt_classifier.predict(X_test)
     11
     12 # Evaluate the model
     13 accuracy_dt = accuracy_score(y_test, y_pred_dt)
     14 classification_report_dt = classification_report(y_test, y_pred_dt)
     15
     16 print(f"Decision Tree Accuracy: {accuracy_dt:.2f}")
     17 print("Decision Tree Classification Report:\n", classification_report_dt)
     18
```



Decision Tree Accuracy: 0.97

Decision Tree Classification Report:

	precision	recall	f1-score	support
acc	0.99	0.90	0.94	83
good	0.62	0.91	0.74	11
unacc	0.99	1.00	1.00	235
vgood	0.82	0.82	0.82	17
accuracy			0.97	346
macro avg	0.86	0.91	0.88	346
weighted avg	0.97	0.97	0.97	346

Decision tree model has same overall accuracy as SVM and outperformed all others in this metric.

The model also shows better performance in all classes with only marginally lower performance in unacc class label than SVM.

7.6 Random Forest

Random Forest is an ensemble method built on decision trees, providing better generalisation and reducing overfitting.

```
[26] 1 from sklearn.ensemble import RandomForestClassifier
      2
      3 # Init of Random Forest
      4 rf_classifier = RandomForestClassifier()
      5
      6 # Train the classifier
      7 rf_classifier.fit(X_train, y_train)
      8
      9 # Make predictions on the test set
     10 y_pred_rf = rf_classifier.predict(X_test)
     11
     12 # Evaluate the model
     13 accuracy_rf = accuracy_score(y_test, y_pred_rf)
     14 classification_report_rf = classification_report(y_test, y_pred_rf)
     15
     16 print(f"Random Forest Accuracy: {accuracy_rf:.2f}")
     17 print("Random Forest Classification Report:\n", classification_report_rf)
     18
```

```
<ipython-input-26-ed49d8c5d23a>:7: DataConversionWarning: A column-vector y
      rf_classifier.fit(X_train, y_train)
Random Forest Accuracy: 0.96
Random Forest Classification Report:
              precision    recall  f1-score   support

     acc           0.97         0.89         0.93         83
    good           0.53         0.82         0.64         11
   unacc           1.00         1.00         1.00        235
   vgood           0.88         0.88         0.88         17

 accuracy                   0.96         346
 macro avg           0.85         0.90         0.86         346
 weighted avg          0.97         0.96         0.96         346
```

Random forest Classifier performs better than KNN, Naïve bayes, Logistic regression but fails to trump SVM and Decision Tree.

8.Conclusion

Based on the metrics used, we have seen that Decision Tree and SVM proved to be better than all others(even Random Forest was outperformed marginally).

Overall Decision Tree Classifier is better than SVM for the following-

- Overall better performance over all class labels in most metrics.
- Performed better in least represented class label.

The poor performance in good class label could be due to its under-representation in the dataset which caused all models to performance distinctly poor in that class's prediction.

Decision trees perform automatic feature selection by selecting the most important features for splitting and deciding which features in the cars attribute proved more required for the class labels improving the accuracy.